

# z0p: A Zero-Dependency Transparent STARK Stack in Safe Rust

nul0

*Future Proofs Tech*

futureproofs@proton.me

September 2026

## Abstract

We present **z0p**, a complete zero-knowledge proving stack implemented from first principles in safe Rust with no dependencies beyond the compiler and the operating system. Binary fields, the BabyBear prime field, BLAKE3, Merkle trees, a parallel NTT, an arity-4 FRI polynomial commitment, and four STARK systems (Fibonacci AIR, range checks, table membership, and hash-folded IVC, plus LogUp lookups and zero-knowledge masking throughout) are all owned by a single crate under `#![forbid(unsafe_code)]`. The system is transparent (no trusted setup) and hash-based, hence plausibly post-quantum. Production parameters target  $\approx 100$  conjectured soundness bits (rate 1/16, 40 queries, 20-bit grinding). Measured on a 16-core machine, a 1,048,576-step computation proves in 6.4 s ( $\approx 6.1 \mu\text{s}/\text{step}$ ), verifies in 0.29 s, with a 160,401-byte proof. Validation combines 155 unit tests with two adversarial suites ( $\approx 800$  hostile attempts per run, zero acceptances, zero panics), wire-exact size assertions on every benchmark, and machine-checked regression gates.

## 1 Introduction

Practical zero-knowledge infrastructure asks its users to trust a great deal: pairing-friendly elliptic curves, multi-party setup ceremonies, deep dependency trees, and provers whose soundness stories live in slide decks. At the same time, large-scale quantum computing threatens every discrete-logarithm and pairing-based proof system in deployment. The systems that survive that transition will be hash-based, transparently set up, small enough to audit, and tested hard enough to believe.

**Contributions.** z0p is, to our knowledge, the first *complete* STARK stack — fields, hash, Merkle trees, NTT, polynomial commitment, AIR systems, and zero-knowledge masking — implemented from scratch in safe Rust with zero dependencies:

- A transparent, plausibly post-quantum design: BLAKE3-256 digests and FRI Reed–Solomon commitments over BabyBear. No curves, no ceremonies.
- Production soundness sizing in one call (`Params::for_security(..., 100)`): blowup 16, queries, grinding, final layer, and cap height jointly optimized for  $\approx 100$  conjectured bits.
- Competitive scalar performance with parallelism over `std::thread` only: 1M-step proofs at  $\approx 163\text{k}$  steps/s, millisecond verification, logarithmically growing proofs.
- Adversarial validation as a first-class artifact: every forgery class we could construct is tested continuously, including a fuzzer that once caught a genuine proof-malleability hole (since closed by binding dimensions into the transcript).

## 1.1 Related work

STARKs [2] and the FRI proximity test [3], strengthened by DEEP-FRI [4] and compiled via BCS-style Fiat–Shamir [1], are our direct lineage. Production STARK toolchains (notably Plonky2/Polygon Zero [8] and Risc Zero) share the BabyBear-plus-FRI engineering core but retain large dependency footprints. We differ in taking dependency-freedom and auditable minimalism as hard constraints, and in publishing adversarial test evidence alongside performance claims. Our lookups follow the LogUp logarithmic-derivative design [7]. The hash is BLAKE3 [6]; Fiat–Shamir is classical [5].

## 2 Preliminaries

**Fields.** Program words use binary fields  $GF(2^8)$ ,  $GF(2^{64})$ ,  $GF(2^{128})$ . Polynomial domains use BabyBear  $\mathbb{F}_p$ ,  $p = 2^{31} - 2^{27} + 1 = 2013265921$  with generator 31;  $p - 1 = 2^{27} \cdot 15$ , so  $2^k$ -roots of unity exist for  $k \leq 27$ . Elements serialize as 4 canonical little-endian bytes; multiplication is Barrett reduction, inversion Fermat ( $0 \mapsto 0$ ), batch inversion Montgomery’s trick.

**Domains.** Trace domain  $H = \{1, \omega, \dots, \omega^{N-1}\}$ ; evaluation domain  $D = \gamma \cdot \langle \Omega \rangle$  of size  $M = N \cdot \beta$  ( $\beta$  the blowup), coset shift  $\gamma = 31$ ;  $M \leq 2^{27}$  enforced everywhere. Vanishing polynomial  $Z_H(x) = x^N - 1$ .

**Hash and Merkle trees.** Unkeyed BLAKE3-256, implemented from the specification and validated against reference vectors including chunk boundaries. Trees use  $\text{leaf}(d) = H(0x00\|d)$ ,  $\text{node}(l, r) = H(0x01\|l\|r)$  (second-preimage resistant). *Caps* expose an upper tree layer in the commitment so openings carry only middle path segments: column trees drop  $\min(h, \text{depth})$  levels, FRI quad layers  $\min(h, \text{depth} - 2)$  (the verifier rebuilds the bottom two levels from revealed values).

**Transcripts.** Append-only, length-prefix-framed Fiat–Shamir transcript with a full domain-separation schedule (FRI: `fri-dims` first, then per-digest `fri-layer-root`, three `fri-fold` challenges per layer, `fri-final`, `fri-pow`, per-query `fri-query`; per-protocol `params/column/alpha` labels; IVC `fold/step`, `fold/trace`, five `fold/alpha` per step plus per-step `fold/gamma`; lookup and sumcheck labels analogously). Proof-of-work grinding finds the smallest nonce with  $d$  leading zero bits under the history commitment; difficulties above 256 always fail.

## 3 FRI polynomial commitment

`commit` encodes  $n$  coefficients to  $M = n\beta$  coset evaluations (shifted NTT); `commit_evals` takes caller-supplied LDE values of declared degree  $\leq \text{bound} < M$ . Both enforce the rate gate  $\text{bound} + 1 \leq M/2$  at prove *and* verify time. Prover-echoed dimensions are absorbed under `fri-dims` before any challenge, closing proof malleability of those bytes.

Layers are bit-reversed; each adjacent quad sits at  $(x, -x, j_0x, -j_0x)$  for the fixed fourth root  $j_0$ , sharing  $x^4$ . Writing  $P(X) = A + BX + CX^2 + DX^3$  in  $Y = X^4$ , each layer folds as  $A + r_0B + r_1C + r_2D$  with fresh challenges, quartering the degree; point inverses are batch-inverted once per layer and folding is multi-threaded. Proofs carry the final layer in the clear plus, per query, all four values of every folded quad with middle path segments. The verifier replays the transcript, checks dimensions, rate, cap shapes, and PoW, verifies the final layer by divided differences against residual bound  $\text{bound} \gg 2f$ , then checks every Merkle path and fold relation, chaining intermediate parents forward.

**Conjecture 1** (FRI soundness sizing). *Under the Johnson-radius heuristic ( $\delta \approx 1 - \sqrt{\rho}$ ), rate  $\rho = 1/8$  with 12 queries gives  $\approx 18$  bits (bench shape); rate  $\rho = 1/16$  with 40 queries plus 20-bit grinding gives  $80 + 20 = 100$  bits (production shape). Caps affect size only.*

## 4 AIR systems

All application protocols share one pattern: interpolate columns, extend on the LDE coset, Merkle-commit joint rows, evaluate a composition vector (geometric vanishing recurrence with Montgomery-batched denominators, bit-identical to scalar evaluation), FRI-commit it, open trace columns at FRI query positions, and re-verify the identity per query.

**Fibonacci AIR.** Trace  $A[0] = a_0$ ,  $B[0] = b_0$ ,  $A[i+1] = B[i]$ ,  $B[i+1] = A[i] + B[i]$ , claim  $B[N-1] = b_{\text{last}}$ . With  $w_{\text{last}} = \omega^{N-1}$  and  $Z_{\text{trans}}(x) = (x^N - 1)/(x - w_{\text{last}})$ :

$$\begin{aligned} C(x) = & \alpha_0 \frac{a' - b}{Z_{\text{trans}}(x)} + \alpha_1 \frac{b' - (a + b)}{Z_{\text{trans}}(x)} + \alpha_2 \frac{a - a_0}{x - 1} \\ & + \alpha_3 \frac{b - b_0}{x - 1} + \alpha_4 \frac{b - b_{\text{last}}}{x - w_{\text{last}}}. \end{aligned} \tag{1}$$

Dividing (not multiplying) by  $Z_{\text{trans}}$  is load-bearing: the quotient is low-degree iff the transition holds on every covered row.

**Range checks.** Value column plus  $n \leq 24$  bit columns: decomposition  $(V - \sum 2^j C_j)/Z_H$ , per-bit boolean quotients  $C_j(C_j - 1)/Z_H$ , and a boundary quotient.

**Table membership.**  $C(x) = \alpha_{\text{member}} \prod_{t \in T} (v - t)/Z_H(x) + \alpha_{\text{bound}}(v - \text{value})/(x - 1)$ ; small tables by design.

**Folding (IVC).** Chained segments share one batched FRI commitment over  $\sum_s \gamma_s C_s$  with per-step challenges; one proof covers the chain.

**Lookups.** LogUp-style: multiplicities,  $\beta$  challenge, quotients  $1/(\beta + w)$ ,  $m/(\beta + t)$ , rational-sum equality and four sumchecks; proving fails closed on out-of-table witnesses.

**Zero knowledge.** Masked proving blinds columns with vanishing multiples from a ChaCha20 stream ( $n_{\text{rand}} \geq$  queries enforced; key material zeroized on drop). Openings become witness-independent; verification is unchanged, and masked false statements still fail.

## 5 Encodings and parameters

Wire format (VERSION = 1): counts as *u32*, dimensions as *u64*, fields as 4 canonical bytes ( $< p$  enforced), digests as 32 bytes, direction bits 0/1. The decoder fails closed on wrong versions, trailing bytes, truncation, non-canonical fields, and counts past hard caps ( $2^{20}$  queries, 64 folds, 128 path steps, 256 cap digests,  $2^{27}$  elements,  $2^{20}$  IVC steps) before any allocation. Validation: power-of-two domains, blowup  $\geq 4$  for AIR,  $M/\text{final}$  a power of four,  $M \leq 2^{27}$ .

## 6 Security analysis

**Assumptions.** BLAKE3 collision/preimage resistance; FRI soundness heuristic; Fiat–Shamir in the (Q)ROM. No discrete-log, pairing, or setup assumptions: Shor-inert by construction.

**Enforced invariants.** Bilateral rate gate; verifier-derived domains and query indices; final divided-difference check at true folded points; cap/path shape pins before hashing; **fri-dims** dimension binding; canonical-only decoding; PoW bound to history with queries sampled after.

**Adversarial validation.** Three layers: 155 unit tests (KATs, Horner cross-checks, parallel-vs-sequential identity); a 13-vector security suite (tampering of every query/fold/value/path/digest, 300 single-bit-flip fuzz trials, PoW forgery, rate gates both directions, statement replay); and a 10-vector red-team suite (cross-proof transplants, false statements, IVC surgery, codec-limit abuse, transcript abuse, exact rate boundaries, masked soundness, and a 500-input garbage storm that must fail closed without panicking). Standing:  $\approx 800$  hostile attempts per run, zero acceptances, zero panics, in debug and release.

## 7 Implementation and performance

Safe Rust, zero dependencies, parallelism over `std::thread` only (NTT stages  $\geq 2^{16}$ , folding, Merkle levels; all bit-identical). Key optimizations: Barrett field arithmetic, table-driven NTT, single-chunk hash fast paths, batched composition (geometric vanishing + Montgomery denominators), verifier on-demand domain points, Merkle caps. Measured 2026-09-26, release, medians, 16-core/31 GB, `rustc` 1.97.1:

| Workload (production, $\approx 100$ bits) | Prove        | Verify        | Size             |
|-------------------------------------------|--------------|---------------|------------------|
| AIR 1,024 steps                           | 23 ms        | 1.4 ms        | 48,221 B         |
| AIR 8,192 steps                           | 75 ms        | 2.8 ms        | 74,809 B         |
| AIR 65,536 steps                          | 370 ms       | 16 ms         | 107,609 B        |
| AIR <b>1,048,576</b> steps                | <b>6.4 s</b> | <b>0.29 s</b> | <b>160,401 B</b> |
| FRI poly-1024                             | 37 ms        | 1.0 ms        | 33,382 B         |

Marginal cost  $\approx 6.1 \mu\text{s}/\text{step}$  ( $\approx 163\text{k}$  steps/s); proof sizes are deterministic and asserted byte-for-byte every run. Primitives: BLAKE3 1 MiB in 1.51 ms, BabyBear 134.6 Mops/s, NTT  $2^{14}$  in 0.36 ms, 64k-leaf Merkle tree in 3.88 ms.

## 8 Limitations and future work

Bench defaults are fast but weak and must not be mistaken for production soundness; large-table lookups await a multilinear PCS; vectorized hashing (pending a **rust-version** policy decision), recursion, and a zkVM are roadmap items; no external audit has yet been performed, for which the suites, gates, versioned wire format, and pinned baselines are the evidence package.

## References

- [1] E. Ben-Sasson, A. Chiesa, and N. Spooner. Interactive oracle proofs. In *TCC 2016*.
- [2] E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev. Scalable, transparent, and post-quantum secure computational integrity. Cryptology ePrint Archive, Report 2018/046, 2018.
- [3] E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev. Fast Reed-Solomon interactive oracle proofs of proximity. In *ICALP 2018*.
- [4] E. Ben-Sasson, D. Bentov, Y. Horesh, and M. Riabzev. DEEP-FRI: Sampling outside the box improves soundness. Cryptology ePrint Archive, Report 2019/336, 2019.

- [5] A. Fiat and A. Shamir. How to prove yourself: practical solutions to identification and signature problems. In *CRYPTO 1986*.
- [6] J.-P. Aumasson, S. Neves, J. O'Connor, and Z. Wilcox-O'Hearn. BLAKE3: one function, fast everywhere. <https://github.com/BLAKE3-team/BLAKE3>, 2020.
- [7] U. Haböck. Multivariate lookups based on logarithmic derivatives. Cryptology ePrint Archive, Report 2022/1530, 2022.
- [8] Polygon Zero Team. Plonky2: fast recursive arguments with PLONK and FRI. <https://github.com/0xPolygonZero/plonky2>, 2022.